


Operating Systems Practice

2: System calls


Kilho Lee

Operating system



System Call

- An interface for accessing the kernel from user space




Trap Handling Process

- Intel architecture



Trap Handling Process (Cont'd)

- Process P can only see its own memory because of user mode (other areas, including kernel, are hidden)




Trap Handling Process (Cont'd)

- Process P wants to call kill () system call



Trap Handling Process (Cont'd)

```
static int(*syscalls[]) (void)      (syscall.c)
```




```
movl $6, %eax;           int $64
```

syscall-table index




Trap Handling Process (Cont'd)

```
struct gatedesc idt[256]    (trap.c)
```



```
movl $6, %eax;           int $64  
                         ^  
                         |  
                         syscall-table index
```

Trap Handling Process (Cont'd)




movl \$6, %eax;
int \$64

syscall-table index

trap-table index

Trap Handling Process (Cont'd)

Kernel mode: we can do anything!



movl \$6, %eax;




syscall-table index

int \$64



trap-table index

Trap Handling Process (Cont'd)



movl \$6, %eax;




syscall-table index

int \$64



trap-table index

Trap Handling Process (Cont'd)



movl \$6, %eax;




syscall-table index

int \$64



trap-table index

Trap Handling Process (Cont'd)



movl \$6, %eax;



syscall-table index

int \$64



trap-table index

Trap Handling Process on Xv6

- user.h

```
4 // system calls
5 int fork(void);
6 int exit(void) __attribute__((noreturn));
7 int wait(void);
8 int pipe(int*);
9 int write(int, void*, int);
10 int read(int, void*, int);
11 int close(int);
12 int kill(int); int kill(int);
13 int exec(char*, char**);
14 int open(char*, int);
15 int mknod(char*, short, short);
16 int unlink(char*);
17 int fstat(int fd, struct stat*);
18 int link(char*, char*);
19 int mkdir(char*);
20 int chdir(char*);
21 int dup(int);
22 int getpid(void);
23 char* sbrk(int);
24 int sleep(int);
25 int uptime(void);
```

Trap Handling Process on Xv6

- usys.S

```

11 SYSCALL(fork)
12 SYSCALL(exit)
13 SYSCALL(wait)
14 SYSCALL(pipe)
15 SYSCALL(read)
16 SYSCALL(write)
17 SYSCALL(close)
18 SYSCALL(kill)
19 SYSCALL(exec)
20 SYSCALL(open)
21 SYSCALL(mknod)
22 SYSCALL(unlink)
23 SYSCALL(fstat)
24 SYSCALL(link)
25 SYSCALL(mkdir)
26 SYSCALL(chdir)
27 SYSCALL(dup)
28 SYSCALL(getpid)
29 SYSCALL(sbrk)
30 SYSCALL(sleep)
31 SYSCALL(uptime)

```

```

1 #include "syscall.h"
2 #include "traps.h"
3
4 #define SYSCALL(name) \
5     .globl name; \
6     name: \
7         movl $SYS_## name, %eax; \
8         int $T_SYSCALL; \
9         ret

```



```

.globl kill;
kill:
    movl $6, %eax;
    int $64;
    ret

```

Trap Handling Process on Xv6

- syscall.h
 - system call number

```
1 // System call numbers
2 #define SYS_fork    1
3 #define SYS_exit    2
4 #define SYS_wait    3
5 #define SYS_pipe    4
6 #define SYS_read    5
7 #define SYS_kill    6
8 #define SYS_exec    7
9 #define SYS_fstat   8
10 #define SYS_chdir   9
11 #define SYS_dup    10
12 #define SYS_getpid 11
13 #define SYS_sbrk   12
14 #define SYS_sleep  13
15 #define SYS_uptime 14
16 #define SYS_open   15
17 #define SYS_write  16
18 #define SYS_mknod  17
19 #define SYS_unlink 18
20 #define SYS_link   19
21 #define SYS_mkdir  20
22 #define SYS_close  21
```

Trap Handling Process on Xv6

- traps.h
 - trap number

```

1 // x86 trap and interrupt constants.
2
3 // Processor-defined:
4 #define T_DIVIDE      0      // divide error
5 #define T_DEBUG       1      // debug exception
6 #define T_NMI        2      // non-maskable interrupt
7 #define T_BRKPT      3      // breakpoint
8 #define T_OFLOW       4      // overflow
9 #define T_BOUND      5      // bounds check
10 #define T_ILLOP      6      // illegal opcode
11 #define T_DEVICE     7      // device not available
12 #define T_DBLFLT    8      // double fault
13 // #define T_COPROC   9      // reserved (not used since 486)
14 #define T_TSS        10     // invalid task switch segment
15 #define T_SEGNP      11     // segment not present
16 #define T_STACK      12     // stack exception
17 #define T_GPFLT     13      // general protection fault
18 #define T_PGFLT     14      // page fault
19 // #define T_RES      15      // reserved
20 #define T_FPERR      16      // floating point error
21 #define T_ALIGN      17      // alignment check
22 #define T_MCHK      18      // machine check
23 #define T SIMDERR   19      // SIMD floating point error
24
25 // These are arbitrarily chosen, but with care not to overlap
26 // processor defined exceptions or interrupt vectors.
27 #define T_SYSCALL    64      // system call
28 #define T_DEFAULT    500     // catchall
29
30 #define T_IRQ0       32      // IRQ 0 corresponds to int T_IRQ
31
32 #define IRQ_TIMER    0
33 #define IRQ_KBD      1
34 #define IRQ_COM1     4
35 #define IRQ_IDE      14
36 #define IRQ_ERROR    19
37 #define IRQ_SPURIOUS 31
38

```

Trap Handling Process on Xv6

- trap.c
 - Interrupt Descriptor Table initialization

```
// Interrupt descriptor table (shared by all CPUs).
struct gatedesc idt[256];
extern uint vectors[]; // in vectors.S: array of 256 entry pointers
struct spinlock tickslock;
uint ticks;

void
tvinit(void)
{
    int i;

    for(i = 0; i < 256; i++)
        SETGATE(idt[i], 0, SEG_KCODE<<3, vectors[i], 0);
    SETGATE(idt[T_SYSCALL], 1, SEG_KCODE<<3, vectors[T_SYSCALL], DPL_USER);

    initlock(&tickslock, "time");
}
```

Trap Handling Process on Xv6

- vectors.S

```
# generated by vectors.pl - do not edit
# handlers
.globl alltraps
.globl vector0
vector0:
    pushl $0
    pushl $0
    jmp alltraps
.globl vector1
vector1:
    pushl $0
    pushl $1
    jmp alltraps
.globl vector2
vector2:
    pushl $0
    pushl $2
    jmp alltraps
    pushl $0
    pushl $64
    jmp alltraps
```

```
# vector table
.data
.globl vectors
vectors:
    .long vector0
    .long vector1
    .long vector2
    .long vector3
    .long vector4
    .long vector5
    .long vector6
    .long vector7
    .long vector8
    .long vector9
    .long vector10
```

Trap Handling Process on Xv6

- trapasm.S

```

# vectors.S sends all traps here.
.globl alltraps
alltraps:
    # Build trap frame.
    pushl %ds
    pushl %es
    pushl %fs
    pushl %gs
    pushal

    # Set up data segments.
    movw $(SEG_KDATA<<3), %ax
    movw %ax, %ds
    movw %ax, %es

    # Call trap(tf), where tf=%esp
    pushl %esp
    call trap
    addl $4, %esp

```




Figure 3-2. The trapframe on the kernel stack

Trap Handling Process on Xv6

- trap.c

```

36 void
37 trap(struct trapframe *tf)
38 {
39     if(tf->trapno == T_SYSCALL) {
40         if(myproc()->killed)
41             exit();
42         myproc()->tf = tf;
43         syscall(); // Line 43
44         if(myproc()->killed)
45             exit();
46         return;
47     }
...

```

```

150 struct trapframe {
151     // registers as pushed
152     uint edi;
153     uint esi;
154     uint ebp;
155     uint oesp;      // user stack
156     uint ebx;
157     uint edx;
158     uint ecx;
159     uint eax;
160
161     // rest of trap frame
162     ushort gs;
163     ushort padding1;
164     ushort fs;
165     ushort padding2;
166     ushort es;
167     ushort padding3;
168     ushort ds;
169     ushort padding4;
170     uint trapno;

```

x86.h

Trap Handling Process on Xv6

- `syscall.c`

```
void
syscall(void)
{
    int num;
    struct proc *curproc = myproc();

    num = curproc->tf->eax;
    if(num > 0 && num < NELEM(syscalls) && syscalls[num] ) {
        curproc->tf->eax = syscalls[num] ();
    } else {
        cprintf("%d %s: unknown sys call %d\n",
                curproc->pid, curproc->name, num);
        curproc->tf->eax = -1;
    }
}
```

```
static int (*syscalls[]) (void) = {  
    [SYS_fork]      sys_fork,  
    [SYS_exit]       sys_exit,  
    [SYS_wait]       sys_wait,  
    [SYS_pipe]       sys_pipe,  
    [SYS_read]       sys_read,  
    [SYS_kill]       sys_kill,  
    [SYS_exec]       sys_exec,  
};  
  
extern int sys_dup(void);  
extern int sys_exec(void);  
extern int sys_exit(void);  
extern int sys_fork(void);  
extern int sys_fstat(void);  
extern int sys_getpid(void);  
extern int sys_kill(void);  
extern int sys_link(void);
```

```
.globl kill;  
kill:  
    movl $6, %eax;  
    int $64;  
    ret
```

Trap Handling Process on Xv6

- sysproc.c

```
29 int
30 sys_kill(void)
31 {
32     int pid;
33
34     if(argint(0, &pid) < 0)
35         return -1;
36     return kill(pid);
37 }
```


- proc.c

```
int
kill(int pid)
{
    struct proc *p;

    acquire(&ptable.lock);
    for(p = ptable.proc; p < &ptable.proc[NPROC]; p++) {
        if(p->pid == pid) {
            p->killed = 1;
```

Trap Handling Process on Xv6

- kill system call



Test with User Program

- Example: kill system call
- user/kill.c

```
1 #include "types.h"
2 #include "stat.h"
3 #include "user.h"
4
5 int
6 main(int argc, char **argv)
7 {
8     int i;
9
10    if(argc < 2){
11        printf(2, "usage: kill pid...\n");
12        exit();
13    }
14    for(i=1; i<argc; i++)
15        kill(atoi(argv[i]));
16    exit();
17 }
```

Test with User Program

- Makefile

```

167 # http://www.gnu.org/software/make/manual,
168 .PRECIOUS: %.o
169
170 U=user
171 UPROGS=\
172     $U/_cat\
173     $U/_echo\
174     $U/_forktest\
175     $U/_grep\
176     $U/_init\
177     $U/_kill\
178     $U/_ln\
179     $U/_ls\
180     $U/_mkdir\
181     $U/_rm\
182     $U/_sh\
183     $U/_stressfs\
184     $U/_usertests\
185     $U/_wc\
186     $U/_zombie\
187     $U/_test_sys\
188
189 fs.img: mkfs README $(UPROGS)
190     ./mkfs fs.img README $(UPROGS)
191

```

- xv6

```

$ ls
.
..
README
cat
echo
forktest
grep
init
kill
ln
ls
mkdir
rm
sh
stressfs
usertests
wc
zombie
console
$


```

(How to) Add user program

- Write your .c code and add it's name to “Makefile”
 - If you write test.c you have to add ‘_test\’ to Makefile.
 - Then, you can execute ‘test’ program on xv6 after booting it

```
UPROGS=\
    _cat\
    _echo\
    _forktest\
    _grep\
    _init\
    _kill\
    _ln\
    _ls\
    _mkdir\
    _rm\
    _sh\
    _stressfs\
    _usertests\
    _wc\
    _zombie\
    _test\
```

References

- Xv6-books
 - <https://pdos.csail.mit.edu/6.828/2018/xv6/book-rev10.pdf>
- Contents
 - **Ch.0: Operating system interface**
 - Ch.1: Operating system organization
 - Ch.2: Page tables
 - Ch.3: Traps, interrupts, and drivers
 - Ch.4: Locking
 - Ch.5: Scheduling
 - Ch.6: File system
 - Ch.7: Summary
 - Appendix A: PC hardware
 - Appendix B: The boot loader

PI. System call Implementation

- Implement your own system calls as follows:
 - getnice
 - setnice
 - ps

setnice / getnice

- A system call that sets/gets a priority of process by using nice value (0~30)
- int **setnice**(int pid, int nice)
 - Set the process priority
 - Argument: process ID, nice value
 - Return value: failure -1 / success 0
- int **getnice**(int pid)
 - Get the process priority
 - Argument: process ID
 - Return value: failure -1 / success process' nice value

ps

- A system call that prints out the process information
 - pid
 - state: RUNNING, RUNNABLE, SLEEPING
 - priority (nice value)
 - runtime (as of now, keep it 0 for future use)

```
$ ps
name      pid      state        priority      runtime      tick 146
init      1        SLEEPING    20            2
sh        2        SLEEPING    20            0
ps        3        RUNNING     20            0
```

Test I

- user/test_nice.c
 - Get nice of init process (must be 5)
 - Get nice of non-existing process (wrong pid) – should return -1
 - Set nice to current process
 - Set nice to non-existing process (wrong pid)
 - should return -1
 - Set wrong nice(<0 or >30) on current process
 - should return -1
 - Get nice of forked process
 - Forked process inherits parent process priority

Test I

- user/test_nice.c

- You should get “OK” messages for all test cases.

```
6 void test_nice()
7 {
8     int pid, nice;
9
10    printf(1, "case 1. get nice value of init process: ");
11    if (getnice(1) == 5)
12        printf(1, "OK\n");
13    else
14        printf(1, "WRONG\n");
15
16    printf(1, "case 2. get nice value of non-existing process: ");
17    if (getnice(100) == -1)
18        printf(1, "OK\n");
19    else
20        printf(1, "WRONG\n");
21
22    printf(1, "case 3. set nice value of current process: ");
23    pid = getpid();
24    setnice(pid, 3);
25    if (getnice(pid) == 3)
26        printf(1, "OK\n");
27    else
28        printf(1, "WRONG\n");
29
30    printf(1, "case 4. set nice value of non-existing process: ");
31    if (setnice(100, 3) == -1)
32        printf(1, "OK\n");
33    else
34        printf(1, "WRONG\n");
35
36    printf(1, "case 5. set wrong nice value of current process: ");
37    if (setnice(pid, -1) == -1 && setnice(pid, 11) == -1)
38        printf(1, "OK\n");
39    else
40        printf(1, "WRONG\n");
41
42    printf(1, "case 6. get nice value of forked process: ");
43    nice = getnice(pid);
44
45    pid = fork();
46    if (pid == 0) { //child
47        if (getnice(getpid()) == nice) {
48            printf(1, "OK\n");
49            exit();
50        }
51        else {
52            printf(1, "WRONG\n");
53            exit();
54        }
55    }
56    else //parent
57        wait();
58 }
```

Test 2

- user/ps.c

- Comment out the 8th line (printf ()) and uncomment the 7th line (ps()), and then run the user program.


```
1 #include "types.h"
2 #include "stat.h"
3 #include "user.h"
4
5 int main(int argc, char **argv)
6 {
7     //ps();
8     printf(1, "ps\n");
9     exit();
10 }
11
```

- You should get the process information as follows (specific numbers may vary):

```
$ ps
name      pid      state        priority        runtime      tick 146
init      1        SLEEPING    20              2
sh        2        SLEEPING    20              0
ps        3        RUNNING     20              0
```


How to Make New System Call on Xv6

- Declare a library routine for your new system call




How to Make New System Call on Xv6

- Define a library routine for your new system call




How to Make New System Call on Xv6

- Update the system call table



How to Make New System Call on Xv6



Hand-in Procedures

- Download template
 - <https://github.com/KilhoLee/xv6-ssu.git> (pull or clone)
 - cp templates/xv6_ssu_syscall.tar.gz [YOUR_DIRECTORY]
 - cd [YOUR_DIRECTORY]
 - tar xvzf xv6_ssu_syscall.tar.gz
 - **** DO CODING ****
- Compress your code with your ID (ex> 20201234)
 - \$tar cvzf xv6_ssu_syscall_20201234.tar.gz xv6_ssu_syscall
 - Must do **\$make clean before compressing**, if not, you will get **penalty**.
 - Double-check whether the source codes are correct!!
- Submit your tar.gz file through **class.ssu.ac.kr**