

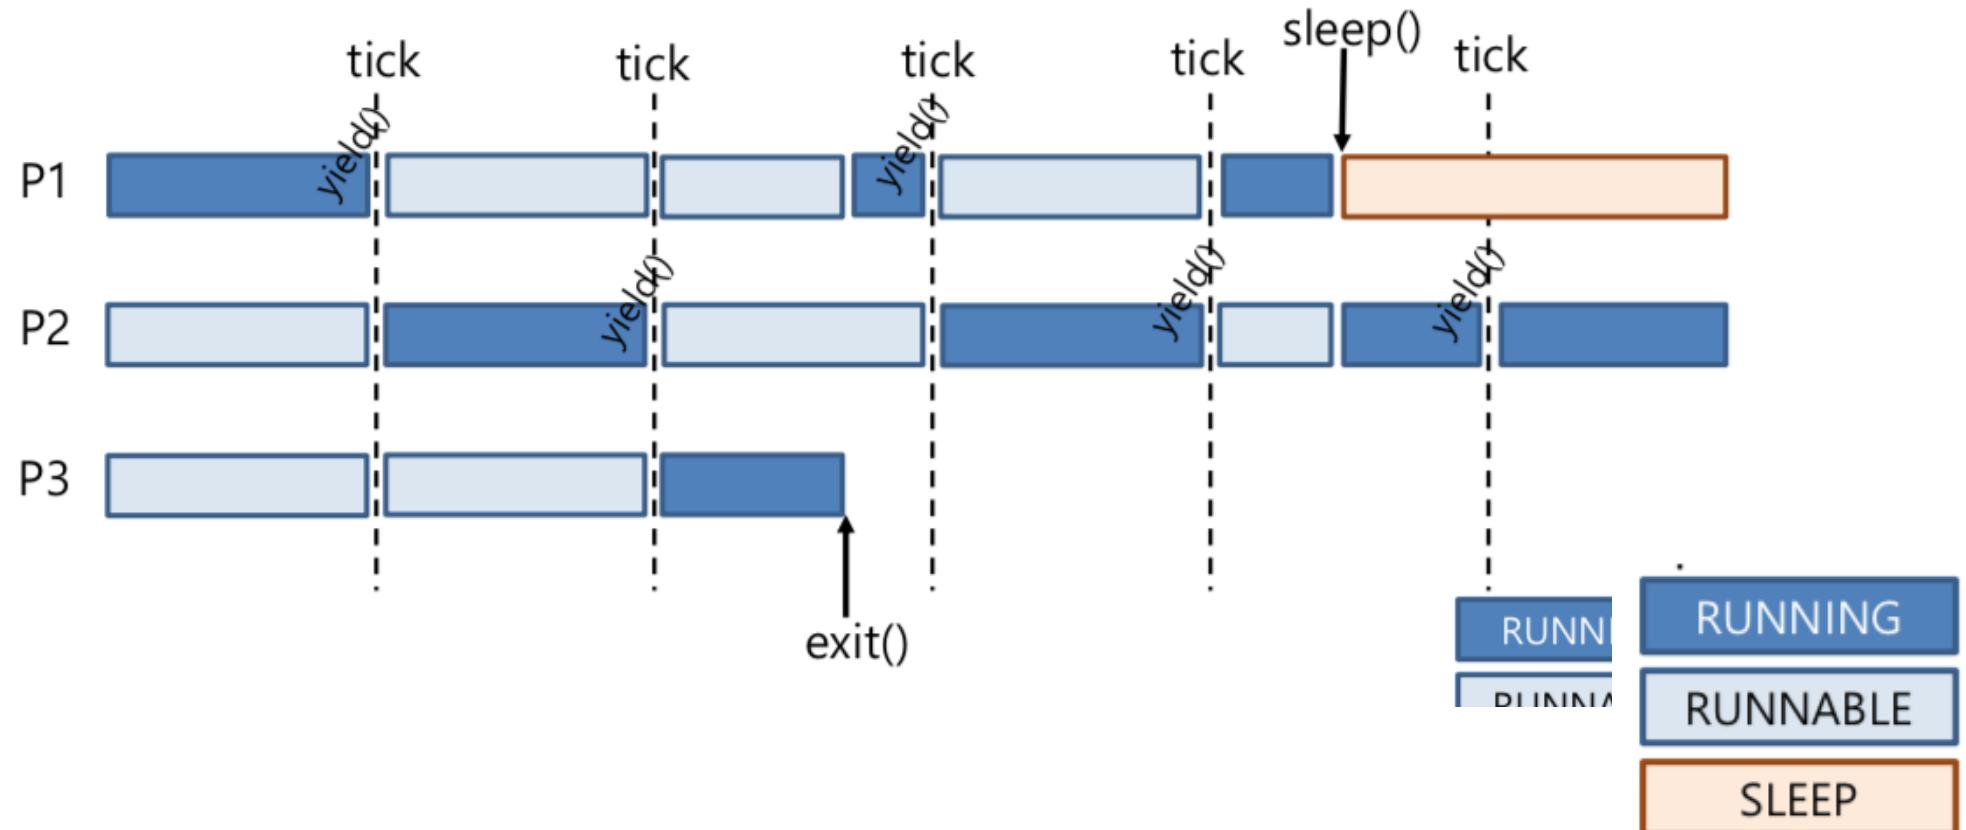
# Operating Systems Practice

3: Scheduler

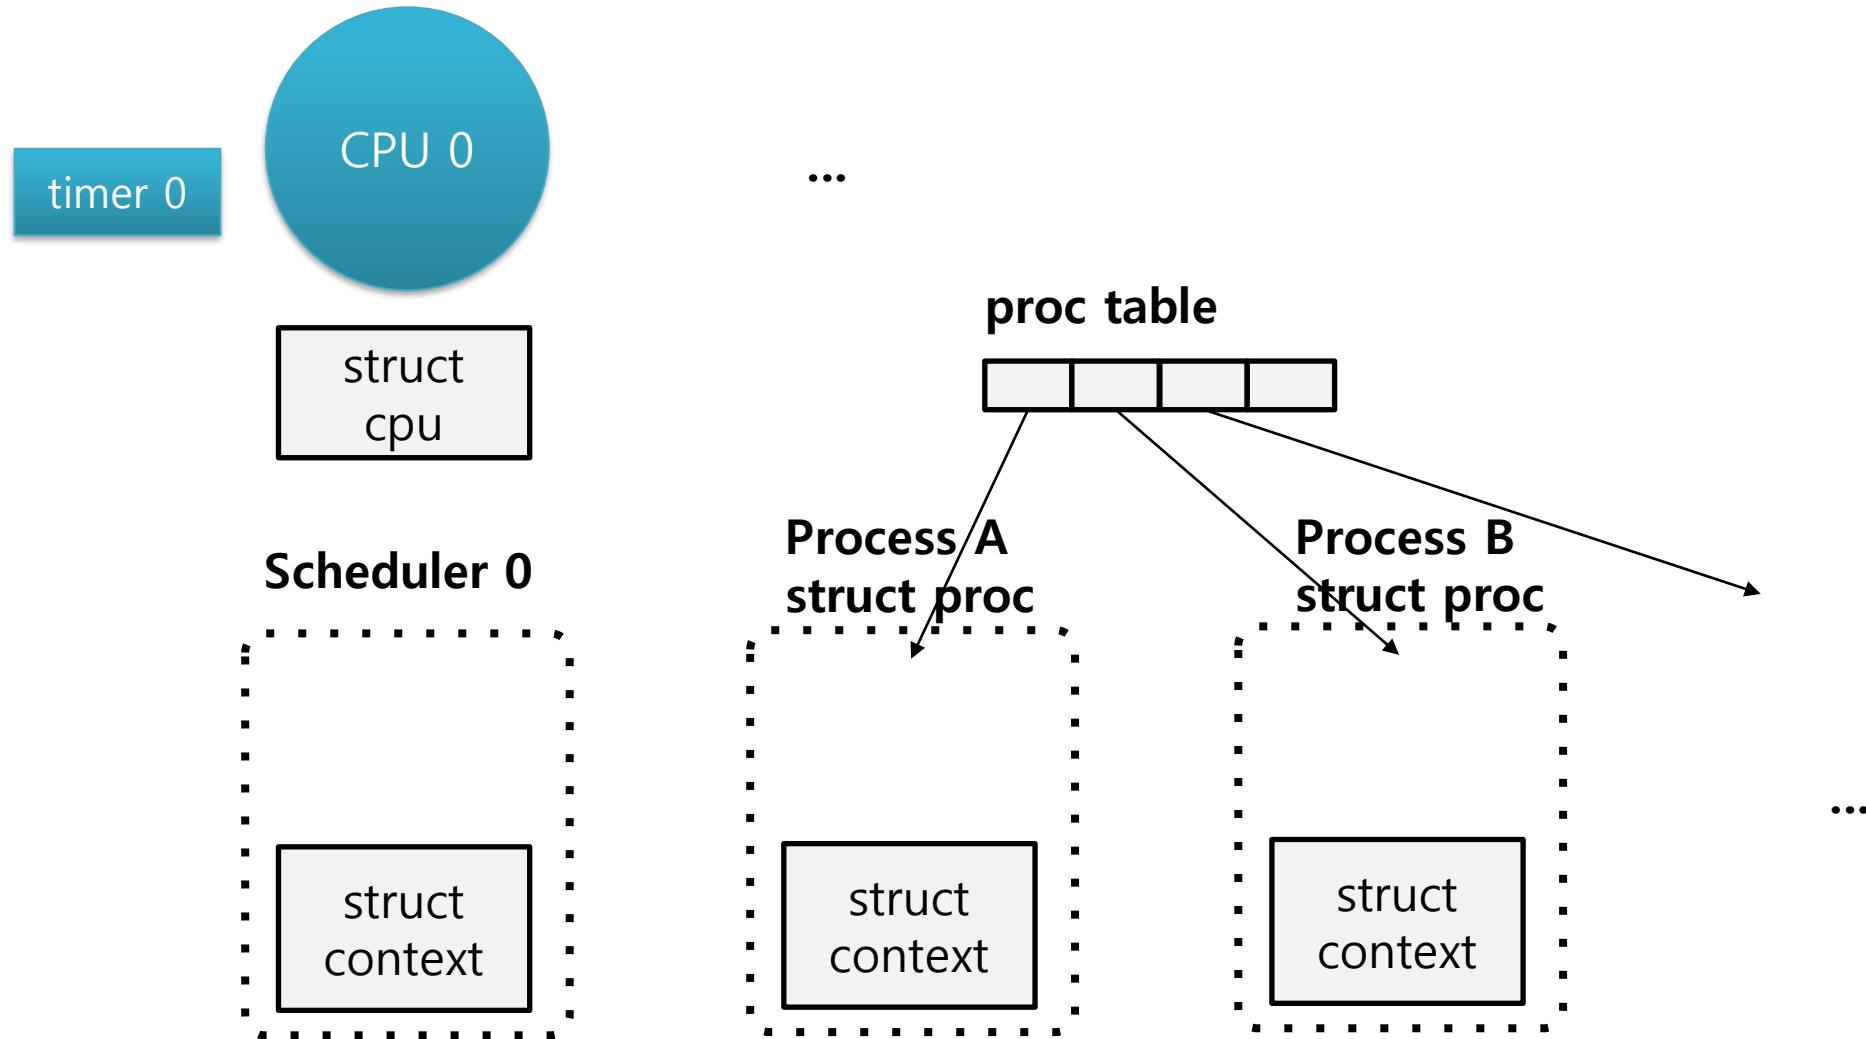
Kilho Lee

# XV6 Scheduler – Round Robin

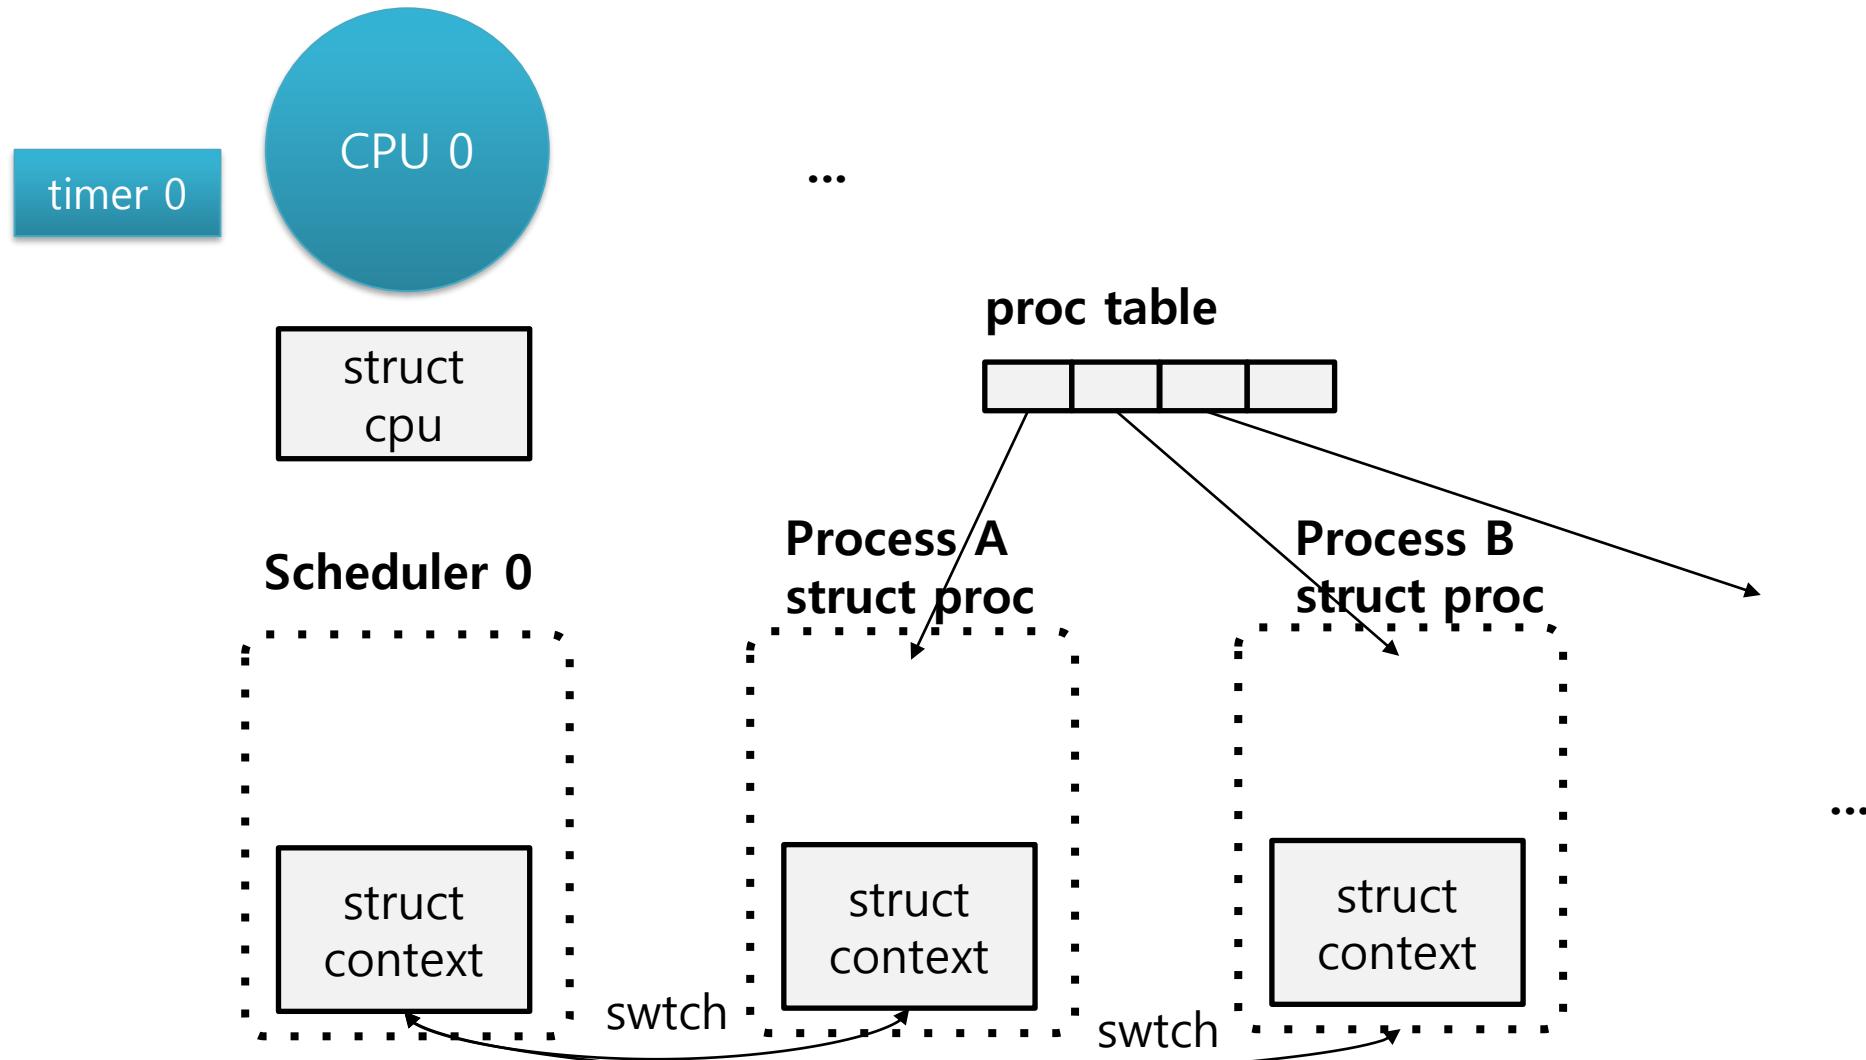
- Timer's interrupt request (IRQ) enforces an yield of a CPU
- A “RUNNABLE” process is chosen to be run in a round-robin manner



# XV6 Data Structures for Scheduling



# XV6 Data Structures for Scheduling



# Process

- proc.h
- procstate
- struct proc

```
35 enum procstate { UNUSED, EMBRYO, SLEEPING, RUNNABLE, RUNNING, ZOMBIE };
36
37 // Per-process state
38 struct proc {
39     uint sz;                                // Size of process memory (bytes)
40     pde_t* pgdir;                           // Page table
41     char *kstack;                           // Bottom of kernel stack for this process
42     enum procstate state;                  // Process state
43     int pid;                               // Process ID
44     struct proc *parent;                   // Parent process
45     struct trapframe *tf;                  // Trap frame for current syscall
46     struct context *context;              // swtch() here to run process
47     void *chan;                            // If non-zero, sleeping on chan
48     int killed;                            // If non-zero, have been killed
49     struct file *ofile[NFILE];            // Open files
50     struct inode *cwd;                   // Current directory
51     char name[16];                        // Process name (debugging)
52     int priority;                          // Process priority
53 };
```

## Proc State

- UNUSED: Not used
- EMBRYO: Newly allocated (not ready for running yet)
- SLEEPING: Waiting for I/O, child process, or time
- RUNNABLE: Ready to run
- RUNNING: Running on CPU
- ZOMBIE: Exited

# Process

- proc.h
- struct context

```
16 //PAGEBREAK: 17
17 // Saved registers for kernel context switches.
18 // Don't need to save all the segment registers (%cs, etc),
19 // because they are constant across kernel contexts.
20 // Don't need to save %eax, %ecx, %edx, because the
21 // x86 convention is that the caller has saved them.
22 // Contexts are stored at the bottom of the stack they
23 // describe; the stack pointer is the address of the context.
24 // The layout of the context matches the layout of the stack in swtch.S
25 // at the "Switch stacks" comment. Switch doesn't save eip explicitly,
26 // but it is on the stack and allocproc() manipulates it.
27 struct context {
28     uint edi;
29     uint esi;
30     uint ebx;
31     uint ebp;
32     uint eip;
33 };
```

# Scheduler

- proc.h
- struct cpu

```
1 // Per-CPU state
2 struct cpu {
3     uchar apicid;           // Local APIC ID
4     struct context *scheduler; // swtch() here to enter scheduler
5     struct taskstate ts;    // Used by x86 to find stack for interrupt
6     struct segdesc gdt[NSEGS]; // x86 global descriptor table
7     volatile uint started;   // Has the CPU started?
8     int ncli;                // Depth of pushcli nesting.
9     int intena;              // Were interrupts enabled before pushcli?
10    struct proc *proc;       // The process running on this cpu or null
11 };
12
13 extern struct cpu cpus[NCPU];
14 extern int ncpu;
15
```

# Scheduler

## main.c

```

17 int
18 main(void)
19 {
20     kinit1(end, P2V(4*1024*1024)); // phys page allocator
21     kvmalloc(); // kernel page table
22     mpinit(); // detect other processors
23     lapicinit(); // interrupt controller
24     seginit(); // segment descriptors
25     picinit(); // disable pic
26     ioapicinit(); // another interrupt controller
27     consoleinit(); // console hardware
28     uartinit(); // serial port
29     pinit(); // process table
30     tvinit(); // trap vectors
31     binit(); // buffer cache
32     fileinit(); // file table
33     ideinit(); // disk
34     startothers(); // start other processors
35     kinit2(P2V(4*1024*1024), P2V(PHYSTOP)); // must come after
36     userinit(); // first user process
37     mpmain(); // finish this processor's setup
38 }
39

```

```

50 // Common CPU setup code.
51 static void
52 mpmain(void)
53 {
54     cprintf("cpu%d: starting %d\n", cpuid(), cpuid());
55     idtinit(); // load idt register
56     xchg(&(mycpu()>started), 1); // tell startothers() we're up
57     scheduler(); // start running processes
58 }

```

## proc.c

```

322 void
323 scheduler(void)
324 {
325     struct proc *p;
326     struct cpu *c = mycpu();
327     c->proc = 0;
328
329     for(;;){
330         // Enable interrupts on this processor.
331         sti();
332
333         // Loop over process table looking for process to run.
334         acquire(&ptable.lock);
335         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
336             if(p->state != RUNNABLE)
337                 continue;
338
339             // Switch to chosen process. It is the process's job
340             // to release ptable.lock and then reacquire it
341             // before jumping back to us.
342             c->proc = p;
343             switchuvvm(p);
344             p->state = RUNNING;
345
346             swtch(&(c->scheduler), p->context);
347             switchkvm();
348
349             // Process is done running for now.
350             // It should have changed its p->state before coming back.
351             c->proc = 0;
352         }
353         release(&ptable.lock);
354     }
355 }
356 }

```

Start to execute chosen process

# When to Schedule

- `exit()`, `sleep()`
- timer interrupt (`yield()`)

trap.c

```
36 void
37 trap(struct trapframe *tf)
38 {
```

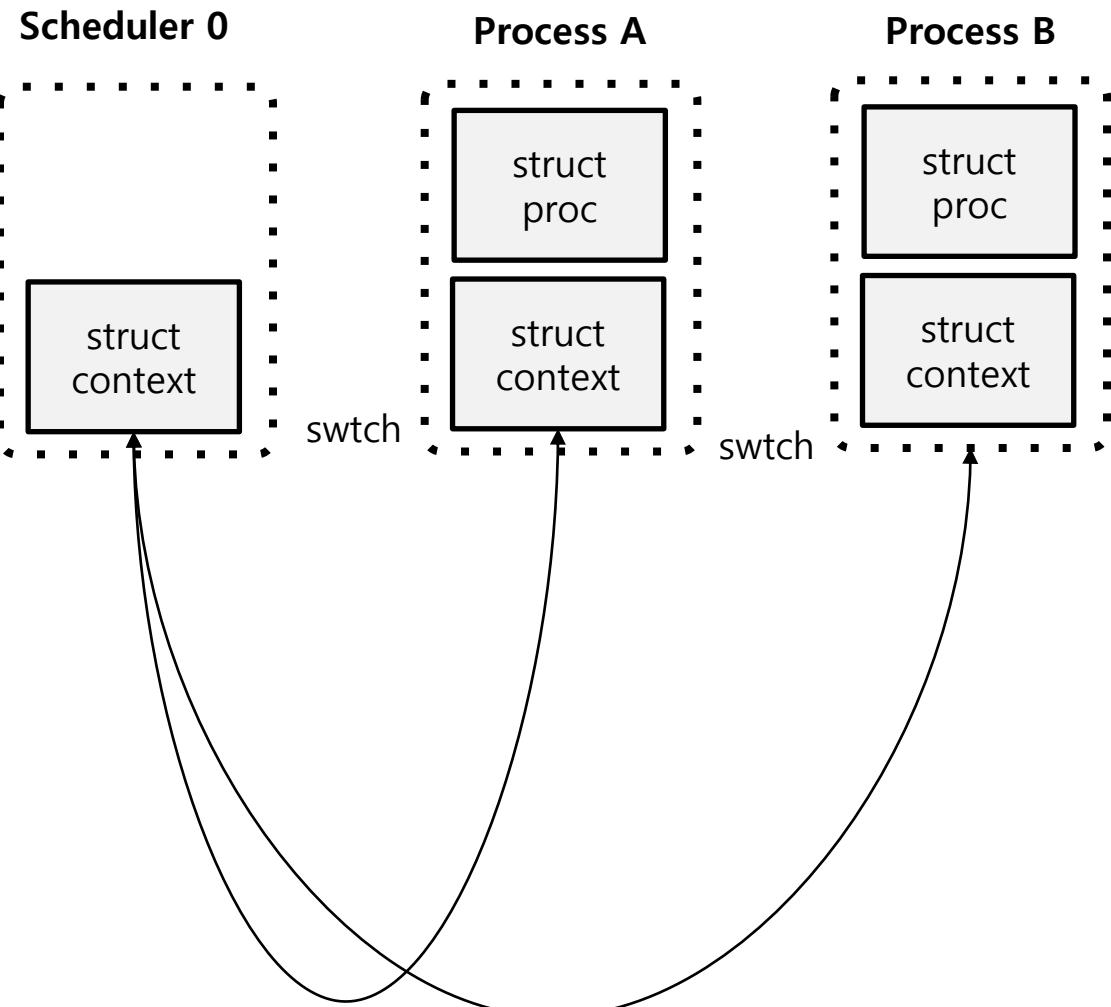
```
103 // Force process to give up CPU on clock tick.
104 // If interrupts were on while locks held, would need to check nlock.
105 if(myproc() && myproc()->state == RUNNING &&
106     tf->trapno == T_IRQ0+IRQ_TIMER)
107     yield();
108 }
```

# How Scheduler works

```

365 void
366 sched(void)
367 {
368     int intena;
369     struct proc *p = myproc();
370
371     if(!holding(&ptable.lock))
372         panic("sched ptable.lock");
373     if(mycpu()>ncli != 1)
374         panic("sched locks");
375     if(p->state == RUNNING)
376         panic("sched running");
377     if(readeflags()&FL_IF)
378         panic("sched interruptible");
379     intena = mycpu()>intena;
380     swtch(&p->context, mycpu()>scheduler);
381     mycpu()>intena = intena;
382 }
383
384 // Give up the CPU for one scheduling round.
385 void
386 yield(void)
387 {
388     acquire(&ptable.lock); //DOC: yieldlock
389     myproc()>state = RUNNABLE;
390     sched();
391     release(&ptable.lock);
392 }

```



# How Scheduler works

```
322 void
323 scheduler(void)
324 {
325     struct proc *p;
326     struct cpu *c = mycpu();
327     c->proc = 0;
328
329     for(;;){
330         // Enable interrupts on this processor.
331         sti();
332
333         // Loop over process table looking for process to run.
334         acquire(&ptable.lock);
335         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
336             if(p->state != RUNNABLE)
337                 continue;
338
339             // Switch to chosen process. It is the process's job
340             // to release ptable.lock and then reacquire it
341             // before jumping back to us.
342             c->proc = p;
343             switchuvvm(p);
344             p->state = RUNNING;
345
346             swtch(&(c->scheduler), p->context);
347             switchkvm();
348
349             // Process is done running for now.
350             // It should have changed its p->state before coming back.
351             c->proc = 0;
352         }
353         release(&ptable.lock);
354     }
355 }
356 }
```

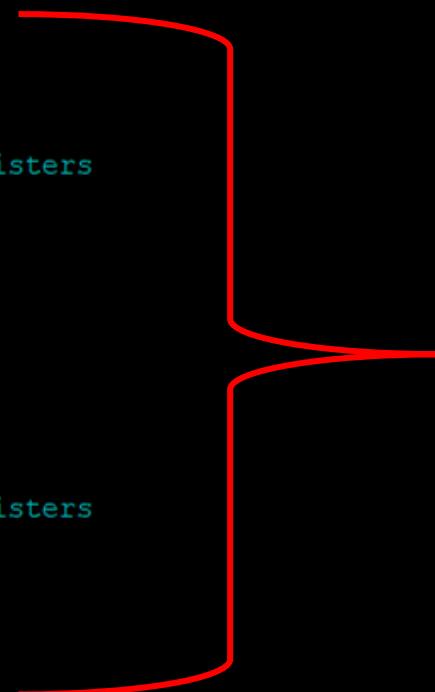


Return from here

# Context switch

- swtch.S

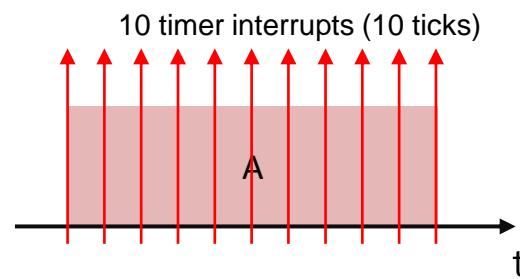
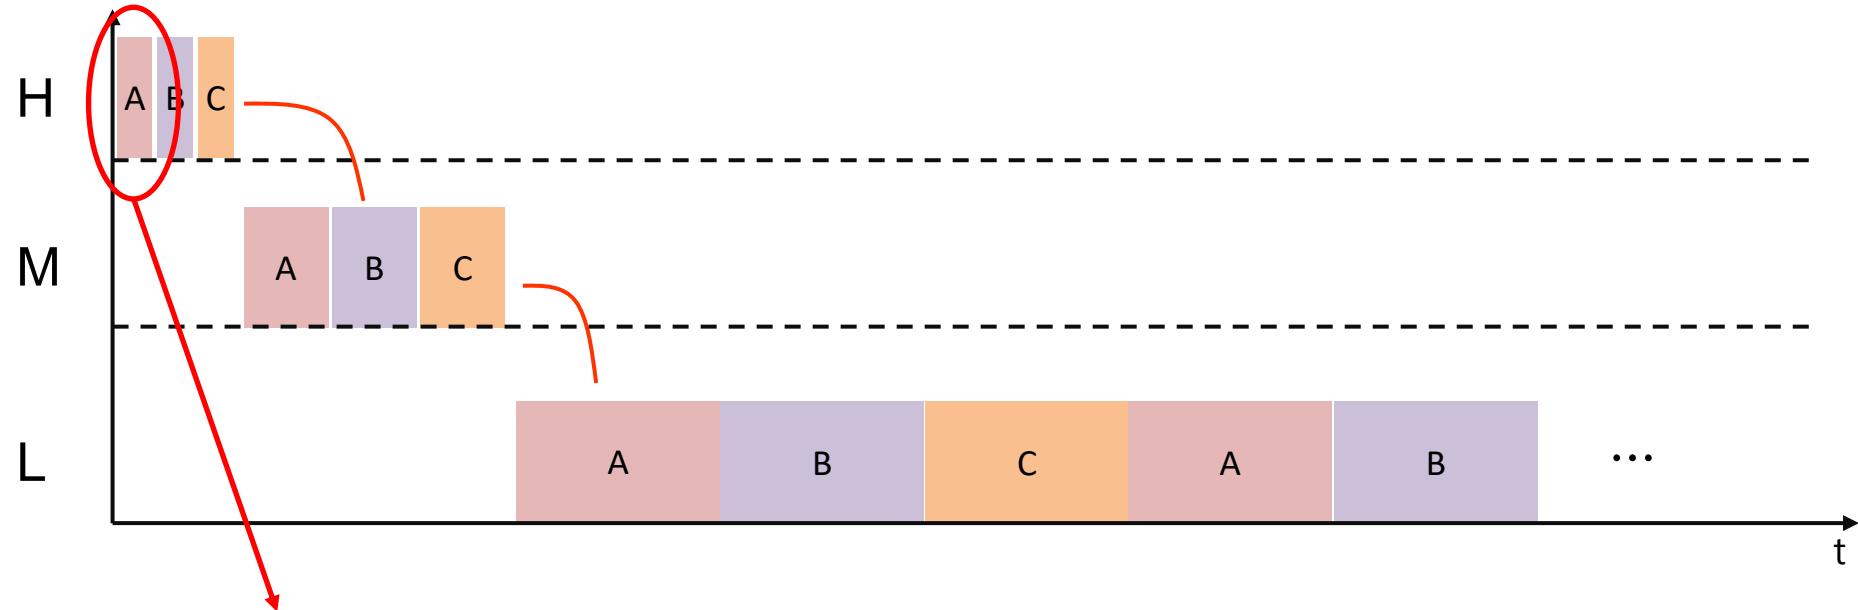
```
1 # Context switch
2 #
3 #     void swtch(struct context **old, struct context *new);
4 #
5 # Save the current registers on the stack, creating
6 # a struct context, and save its address in *old.
7 # Switch stacks to new and pop previously-saved registers.
8
9 .globl swtch
10 swtch:
11     movl 4(%esp), %eax
12     movl 8(%esp), %edx
13
14     # Save old callee-saved registers
15     pushl %ebp
16     pushl %ebx
17     pushl %esi
18     pushl %edi
19
20     # Switch stacks
21     movl %esp, (%eax)
22     movl %edx, %esp
23
24     # Load new callee-saved registers
25     popl %edi
26     popl %esi
27     popl %ebx
28     popl %ebp
29     ret
```



# Project 2. Simple MLFQ Scheduling

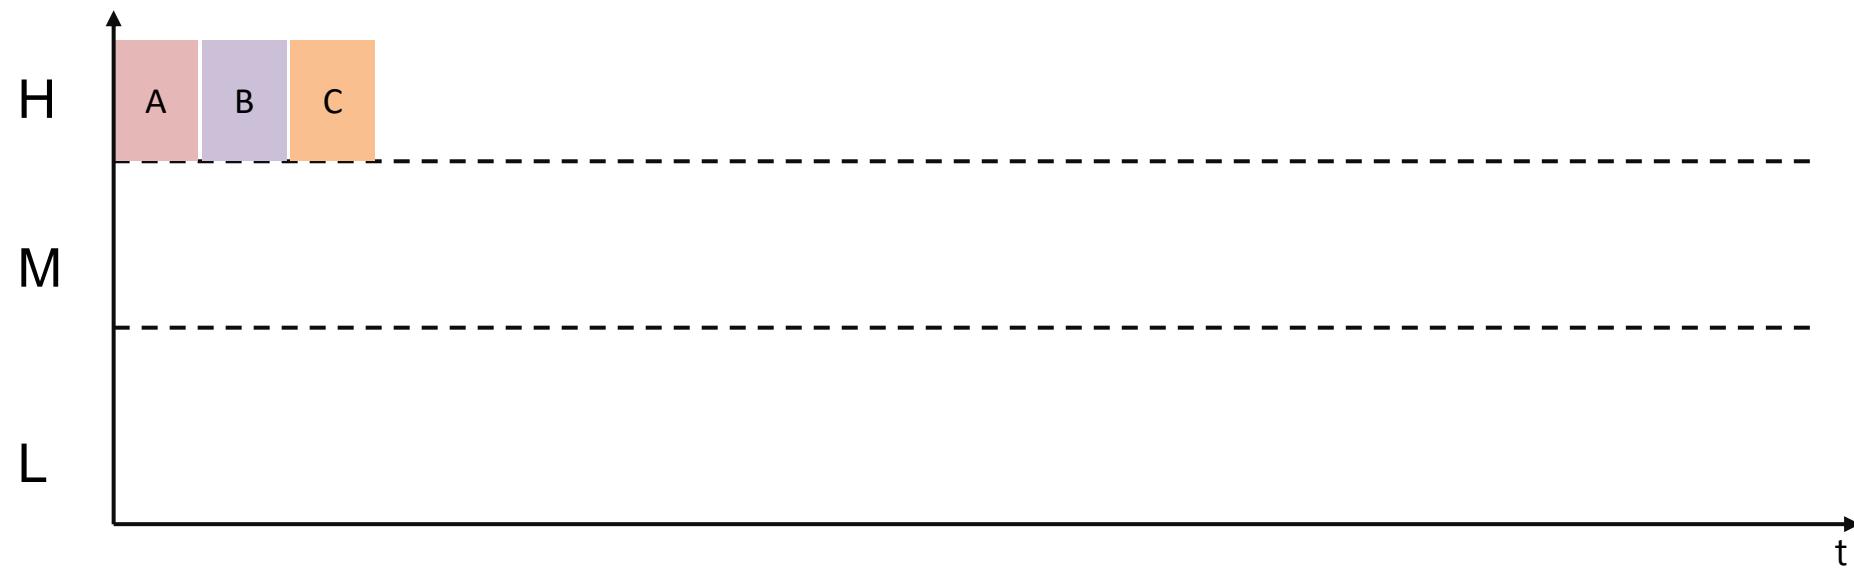
- Implement a MLFQ scheduler that employs the following rules:
  - If priority (process A) > priority (process B), then run the process A.
  - If priority (process A) == priority (process B), then run the process A and B in the RR manner.
  - It has three queues with different priority levels (HIGH(H)/ MID(M)/ LOW(L)).
    - H, M, L queues have the time slices of **10, 20, and 30 timer ticks**.
  - A newly arriving process initially belongs to the H queue.
  - A process which consumes all the time slice moves to the next priority queue.
    - e.g., H → M, M → L
  - If a process gives up the CPU before the time slice is up, it stays at the same priority queue.
- Note
  - No priority boost.
  - No cumulative accounting.

# Example



## test\_rr.c

- Add test\_rr.c as user program
- ./test\_rr
- Each process will consumes 6-7 time ticks (it will reside in the H queue)



# test\_mlfq.c

- 출력은 실행환경의 tick 호출 주기에 따라 약간씩 다를 수 있음.
- Xv6의 기본 round-robin scheduler 가 아닌, time slice 를 모두 사용한 이후에 context switch 되도록 구현
  - HIGH 큐 기준 타임 슬라이스는 10 timer ticks.
  - 테스트 케이스의 P1, P2, P3는 약 5~7 ticks 를 사용하도록 구현됨.
  - 실행 시, P1 → P2 → P3 순서로 실행되어야 함.

```

== TEST START ==
P1 ARRIVED
P1 (high), i = 0, dummy = C0000000
P1 (high), i = 1, dummy = E0000000
P1 (high), i = 2, dummy = 60000000
P1 (high), i = 3, dummy = E0000000
P1 (high), i = 4, dummy = 0
P1 (high), i = 5, dummy = 20000000
P1 (high), i = 6, dummy = 20000000
P1 (high), i = 7, dummy = C0000000
P1 (high), i = 8, dummy = 60000000
P1 (high), i = 9, dummy = 0
P1 (high), i = 10, dummy = A0000000
P1 (high), i = 11, dummy = 40000000
P1 (high), i = 12, dummy = E0000000
P1 (high), i = 13, dummy = 40000000
P1 (high), i = 14, dummy = 0
P1 (high), i = 15, dummy = C0000000
P1 (high), i = 16, dummy = 80000000
P1 (high), i = 17, dummy = 40000000
P1 (high), i = 18, dummy = 0
P1 (high), i = 19, dummy = C0000000
P1 RELEASED
P2 ARRIVED
P2 (high), i = 0, dummy = C0000000
P2 (high), i = 1, dummy = E0000000
P2 (high), i = 2, dummy = 60000000
P2 (high), i = 3, dummy = E0000000
P2 (high), i = 4, dummy = 0
P2 (high), i = 5, dummy = 20000000
P2 (high), i = 6, dummy = 20000000
P2 (high), i = 7, dummy = C0000000
P2 (high), i = 8, dummy = 60000000
P2 (high), i = 9, dummy = 0
P2 (high), i = 10, dummy = A0000000
P2 (high), i = 11, dummy = 40000000
P2 (high), i = 12, dummy = E0000000
P2 (high), i = 13, dummy = 40000000
P2 (high), i = 14, dummy = 0
P2 (high), i = 15, dummy = C0000000
P2 (high), i = 16, dummy = 80000000
P2 (high), i = 17, dummy = 40000000
P2 (high), i = 18, dummy = 0
P2 (high), i = 19, dummy = C0000000
P2 RELEASED
== TEST DONE ==

```

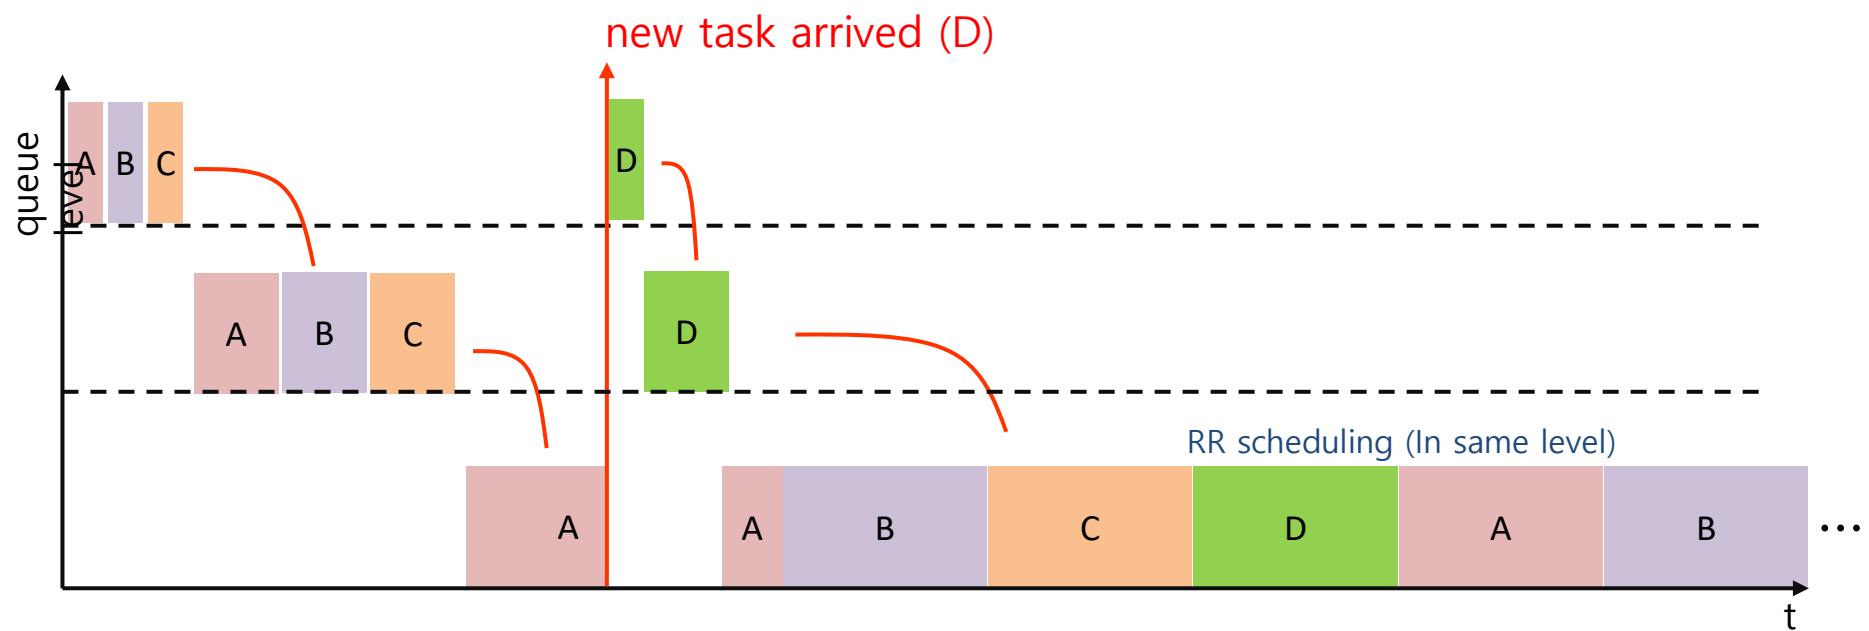
```

P3 ARRIVED
P3 (high), i = 0, dummy = C0000000
P3 (high), i = 1, dummy = E0000000
P3 (high), i = 2, dummy = 60000000
P3 (high), i = 3, dummy = E0000000
P3 (high), i = 4, dummy = 0
P3 (high), i = 5, dummy = 20000000
P3 (high), i = 6, dummy = 20000000
P3 (high), i = 7, dummy = C0000000
P3 (high), i = 8, dummy = 60000000
P3 (high), i = 9, dummy = 0
P3 (high), i = 10, dummy = A0000000
P3 (high), i = 11, dummy = 40000000
P3 (high), i = 12, dummy = E0000000
P3 (high), i = 13, dummy = 40000000
P3 (high), i = 14, dummy = 0
P3 (high), i = 15, dummy = C0000000
P3 (high), i = 16, dummy = 80000000
P3 (high), i = 17, dummy = 40000000
P3 (high), i = 18, dummy = 0
P3 (high), i = 19, dummy = C0000000
P3 RELEASED
== TEST DONE ==

```

# test\_mlfq.c

- Add test\_mlfq.c as user program
- PI output may be different
- ./test\_mlfq



# test\_mlfq.c

- 출력은 실행환경의 tick 호출 주기에 따라 약간씩 다를 수 있음.
- 같은 우선순위 큐 내에서 RR 스케줄링 된다면 문제 없음.
  - 각 우선순위 큐에서 타임 슬라이스마다 프로세스들이 비슷한 반복 횟수를 보이면 문제 없음.  
(예: P1/P4: 3번, P2/P3: 4번 ↗ ok)
  - HIGH, MID 큐의 경우 타임 슬라이스가 작아 프로세스의 남은 작업이 대부분 LOW 큐에서 동작함.
  - LOW 큐의 타임 슬라이스가 크지만 남은 작업을 한 타임 슬라이스 내에 모두 실행하지는 못함.  
(P1, P2, P3, P4가 타임 슬라이스를 재할당 받고 RR scheduling 됨.)
  - 애매한 경우 TA에게 문의

P1, P2, P3 in HIGH queue → RR scheduling.  
Each process exhausts its own time slice(10).

P1, P2, P3 in MID queue → RR scheduling.  
Each process exhausts its own time slice(20).

P1, P2, P3 in LOW queue → RR scheduling.

P4 arrived → P4 in HIGH queue

P4 in MID queue

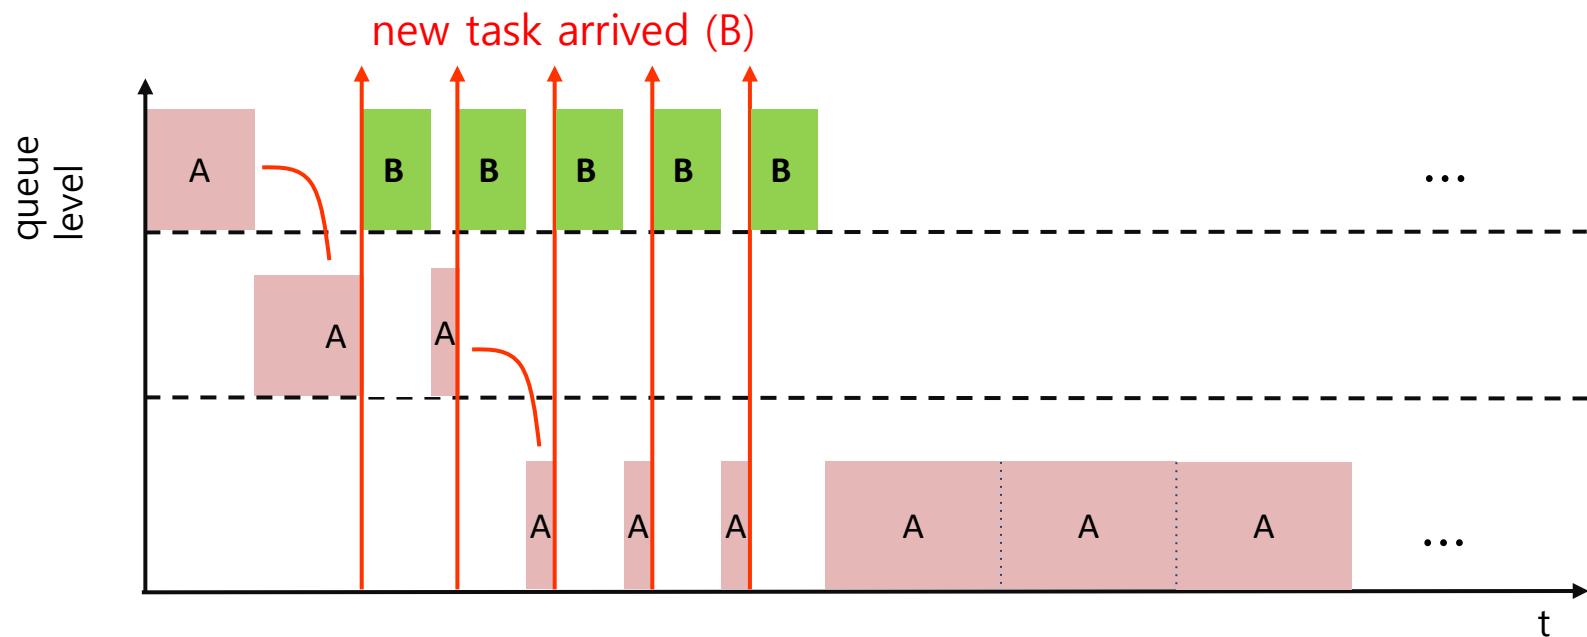
P1, P2, P3, P4 in LOW queue → RR scheduling.  
Each process exhausts its own time slice(30)  
and reset time slice(30).

```
$ test_mlfq
== TEST START ==
P1 (high), i = 0, dummy = 80000000
P2 (high), i = 0, dummy = 80000000
P3 (high), i = 0, dummy = 80000000
P1 (mid), i = 1, dummy = 20000000
P1 (mid), i = 2, dummy = C0000000
P1 (mid), i = 3, dummy = 40000000
P2 (mid), i = 1, dummy = 20000000
P2 (mid), i = 2, dummy = C0000000
P2 (mid), i = 3, dummy = 40000000
P3 (mid), i = 1, dummy = 20000000
P3 (mid), i = 2, dummy = C0000000
P3 (mid), i = 3, dummy = 40000000
P1 (low), i = 4, dummy = C0000000
P1 (low), i = 5, dummy = E0000000
P1 (low), i = 6, dummy = A0000000
P1 (low), i = 7, dummy = 60000000
P1 (low), i = 8, dummy = 20000000
P1 (low), i = 9, dummy = E0000000
P4 ARRIVED
P4 (high), i = 0, dummy = 80000000
P4 (mid), i = 1, dummy = 20000000
P4 (mid), i = 2, dummy = C0000000
P4 (mid), i = 3, dummy = 40000000
P1 (low), i = 10, dummy = 80000000
P2 (low), i = 4, dummy = C0000000
P2 (low), i = 5, dummy = E0000000
P2 (low), i = 6, dummy = A0000000
P2 (low), i = 7, dummy = 60000000
P2 (low), i = 8, dummy = 20000000
P2 (low), i = 9, dummy = E0000000
P2 (low), i = 10, dummy = A0000000
P3 (low), i = 4, dummy = C0000000
P3 (low), i = 5, dummy = E0000000
P3 (low), i = 6, dummy = A0000000
P3 (low), i = 7, dummy = 60000000
P3 (low), i = 8, dummy = 20000000
P3 (low), i = 9, dummy = E0000000
P3 (low), i = 10, dummy = A0000000
P4 (low), i = 4, dummy = C0000000
P4 (low), i = 5, dummy = E0000000
P4 (low), i = 6, dummy = A0000000
P4 (low), i = 7, dummy = 60000000
P4 (low), i = 8, dummy = 20000000
P4 (low), i = 9, dummy = E0000000
P4 (low), i = 10, dummy = A0000000
== TEST DONE ==
```

```
P1 (low), i = 11, dummy = 20000000
P1 (low), i = 12, dummy = C0000000
P1 (low), i = 13, dummy = 40000000
P1 (low), i = 14, dummy = C0000000
P1 (low), i = 15, dummy = E0000000
P1 (low), i = 16, dummy = A0000000
P2 (low), i = 11, dummy = 40000000
P2 (low), i = 12, dummy = 80000000
P2 (low), i = 13, dummy = C0000000
P2 (low), i = 14, dummy = 0
P2 (low), i = 15, dummy = 40000000
P2 (low), i = 16, dummy = 80000000
P3 (low), i = 11, dummy = 40000000
P3 (low), i = 12, dummy = 80000000
P3 (low), i = 13, dummy = C0000000
P3 (low), i = 14, dummy = 0
P3 (low), i = 15, dummy = 40000000
P3 (low), i = 16, dummy = 80000000
P4 (low), i = 11, dummy = 40000000
P4 (low), i = 12, dummy = 80000000
P4 (low), i = 13, dummy = C0000000
P4 (low), i = 14, dummy = 0
P4 (low), i = 15, dummy = 40000000
P4 (low), i = 16, dummy = 80000000
P4 (low), i = 17, dummy = C0000000
P1 (low), i = 17, dummy = 60000000
P1 (low), i = 18, dummy = 20000000
P1 (low), i = 19, dummy = E0000000
P2 (low), i = 17, dummy = C0000000
P2 (low), i = 18, dummy = 0
P2 (low), i = 19, dummy = 40000000
P3 (low), i = 17, dummy = C0000000
P3 (low), i = 18, dummy = 0
P3 (low), i = 19, dummy = 40000000
P4 (low), i = 18, dummy = 0
P4 (low), i = 19, dummy = 40000000
== TEST DONE ==
```

# test\_mlfq2.c

- Add test\_mlfq.c as user program
- PI output may be different
- ./test\_mlfq2



# test\_mlfq2.c

- 출력은 실행환경의 tick 호출 주기에 따라 다를 수 있다.
- HIGH, MID, LOW time slice 는 각각 10, 20, 30으로 설정.
- P1 process 실행 중, P2 process가 주기적으로 생성되어 HIGH time slice 내로 작업 완료 후 종료되어, P1의 starvation 현상이 발생하는 양상이라면 문제 없음.
  - P2 사이에 P1의 실행 횟수는 tick 호출에 따라 달라질 수 있음.
  - Priority Boosting의 필요성을 보여주는 testcase

```

==== TEST START ====
P1 (high), i = 0, dummy = 80000000
P1 (high), i = 1, dummy = 20000000
P1 (mid), i = 2, dummy = C0000000
P1 (mid), i = 3, dummy = 40000000
P2 ARRIVED
P2 (high), i = 0:19, dummy = C0000000
P2 RELEASED
P1 (mid), i = 4, dummy = C0000000
P1 (mid), i = 5, dummy = E0000000
P2 ARRIVED
P2 (high), i = 0:19, dummy = C0000000
P2 RELEASED
P1 (mid), i = 6, dummy = A0000000
P2 ARRIVED
P2 (high), i = 0:19, dummy = C0000000
P2 RELEASED
P1 (mid), i = 7, dummy = 60000000
P1 (low), i = 8, dummy = 20000000
P2 ARRIVED
P2 (high), i = 0:19, dummy = C0000000
P2 RELEASED
P1 (low), i = 9, dummy = E0000000
P2 ARRIVED
P2 (high), i = 0:19, dummy = C0000000
P2 RELEASED
P1 (low), i = 10, dummy = A0000000
P1 (low), i = 11, dummy = 40000000
P1 (low), i = 12, dummy = 80000000
P1 (low), i = 13, dummy = C0000000
P1 (low), i = 14, dummy = 0
P1 (low), i = 15, dummy = 40000000
P1 (low), i = 16, dummy = 80000000
P1 (low), i = 17, dummy = C0000000
P1 (low), i = 18, dummy = 0
P1 (low), i = 19, dummy = 40000000
P1 (low), i = 20, dummy = 80000000
P1 (low), i = 21, dummy = C0000000
P1 (low), i = 22, dummy = 0
P1 (low), i = 23, dummy = 40000000
P1 (low), i = 24, dummy = E0000000
P1 (low), i = 25, dummy = 80000000
P1 (low), i = 26, dummy = 20000000
P1 (low), i = 27, dummy = C0000000
P1 (low), i = 28, dummy = 60000000
P1 (low), i = 29, dummy = 0
==== TEST DONE ====

```

**P1 in HIGH queue.**

**P1 consumes its own time slice(10).**

**P1 in MID queue.**

**P2 arrives, it runs in HIGH queue.**

**P2 is done before it consumes the time slice of HIGH queue.**

**P1 runs in MID queue until it consumes all the time slices (20ticks)**

**P2 arrives repeatedly, it runs in HIGH queue.**

**P1 runs in LOW queue**

# Hand-in Procedures (1/2)

- Download template
  - <https://github.com/KilhoLee/xv6-ssu.git> (pull or clone)
  - tar xvzf xv6\_ssu\_mlfq.tar.gz
- Add **test\_\*.c** to your codes and modify Makefile properly
  - test\_rr.c, test\_mlfq.c, test\_mlfq2.c
- Build with CPUS=1 flag
  - Makefile

```
ifndef CPUS
CPUS := 1
endif
```

# Hand-in Procedures (2/2)

- Compress your code (ID: 20221234)
  - \$tar cvzf xv6\_ssu\_mlfq\_20221234.tar.gz xv6\_ssu\_mlfq
  - Please command \$make clean before compressing
- Submit your tar.gz file through **class.ssu.ac.kr**