

Operating Systems Practice

4: Slab allocator

Kilho Lee

Xv6 Memory Allocator

- `kalloc()` in `kalloc.c`
 - `freelist`: singly linked list
 - Allocation unit: **4KB page**

```
79 // Allocate one 4096-byte page of physical memory.
80 // Returns a pointer that the kernel can use.
81 // Returns 0 if the memory cannot be allocated.
82 char*
83 kalloc(void)
84 {
85     struct run *r;
86
87     if(kmem.use_lock)
88         acquire(&kmem.lock);
89     r = kmem.freelist;
90     if(r)
91         kmem.freelist = r->next;
92     if(kmem.use_lock)
93         release(&kmem.lock);
94     return (char*)r;
95 }
96 
```

```
20 struct {
21     struct spinlock lock;
22     int use_lock;
23     struct run *freelist;
24 } kmem;
```

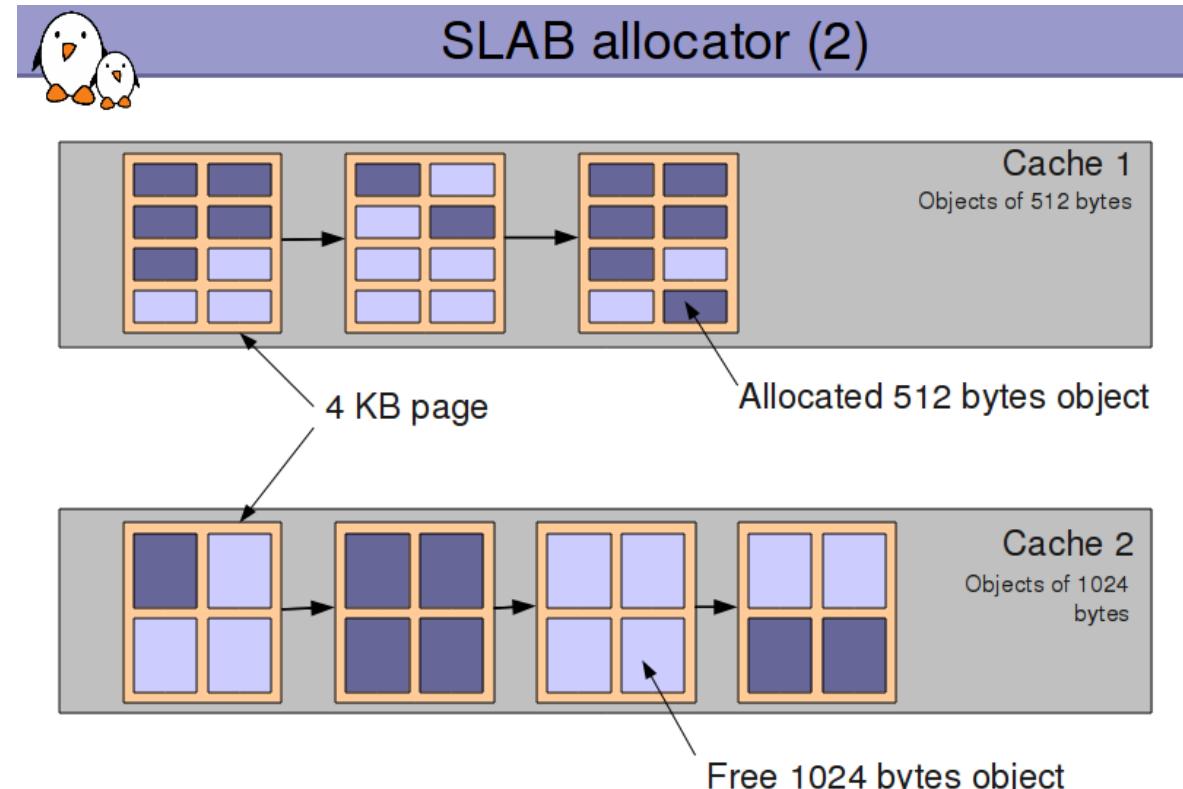
Xv6 Memory Allocator

- kfree in kalloc.c
 - The free page list is embedded into the empty page

```
54 //PAGEBREAK: 21
55 // Free the page of physical memory pointed at by v,
56 // which normally should have been returned by a
57 // call to kalloc(). (The exception is when
58 // initializing the allocator; see kinit above.)
59 void
60 kfree(char *v)
61 {
62     struct run *r;
63
64     if((uint)v % PGSIZE || v < end || V2P(v) >= PHYSTOP)
65         panic("kfree");
66
67     // Fill with junk to catch dangling refs.
68     memset(v, 1, PGSIZE);
69
70     if(kmem.use_lock)
71         acquire(&kmem.lock);
72     r = (struct run*)v;
73     r->next = kmem.freelist;
74     kmem.freelist = r;
75     if(kmem.use_lock)
76         release(&kmem.lock);
77 }
```

Slab Allocator

- A cache that contains a set of objects of the same size
- Used for data structures that are present in many instances in the kernel
 - struct dentry
 - struct inode
 - struct page
 - struct task_struct
 - struct file
 - etc.
- Pre-defined fixed size cache is called general cache
 - 32, 128, 512, ... bytes



Project 3. Byte-level Memory Allocator

- Implement a byte-level slab memory allocator on xv6
- Individual allocator per fixed size of bytes
 - 16, 32, 64, 128, 256, 512, 1024, and 2048 bytes
- New files
 - slab.c, slab.h, sysslab.c, slabtest.c
- Your target
 - Edit the **slab.c** file only

Slab Cache Overview

```
main.c
16 // doing some setup required for memory allocator to work
17 int
18 main(void)
19 {
20     kinit1(end, P2V(4*1024*1024)); // phys page allocator
21     kvmalloc(); // kernel page table
22     mpinit(); // detect other processors
23     lapicinit(); // interrupt controller
24     seginit(); // segment descriptors
25     picinit(); // disable pic
26     ioapicinit(); // another interrupt controller
27     consoleinit(); // console hardware
28     uartinit(); // serial port
29     pinit(); // process table
30     tvinit(); // trap vectors
31     binit(); // buffer cache
32     fileinit(); // file table
33     ideinit(); // disk
34     startothers(); // start other processors
35     kinit2(P2V(4*1024*1024), P2V(PHYSTOP)); // must come after startothers()
)
36     userinit(); // first user process
37     mpmain(); // finish this processor's setup
38 }
```

It requires some initialization in the “main” function
(already added)

Fill up here with your own codes

Project 3. Byte-level Memory Allocator

- slab.h
 - `size`: size of a slab allocator
 - `num_pages`: number of pages of a slab allocator
 - `num_free_objects`: number of free objects of a slab allocator
 - `num_used_objects`: number of used objects of a slab allocator
 - `bitmap`: allocation bitmap for a slab allocator (use a page as a bitmap)
 - `page`: page array for a slab allocator

```
1 #define NSLAB 8 // {16, 32, 64, 128, 256, 512, 1024, 2048}
2 #define MAX_PAGES_PER_SLAB 100
3
4 struct slab {
5     int size;
6     int num_pages;
7     int num_free_objects;
8     int num_used_objects;
9     int num_objects_per_page;
10    char *bitmap;
11    char *page[MAX_PAGES_PER_SLAB];
12 };
13
14 void slabdump();
```

~

DO NOT CHANGE

Project 3. Byte-level Memory Allocator

- slab.c (YOUR TARGET)
 - stable: pre-defined slab allocators & lock
 - void slabinit()
 - Initialization function
 - Allocate one page for bitmap
 - Allocate one page for slab cache
 - char *kmalloc(int size)
 - Byte-level allocation function
 - 0 < size <= 2048
 - Allocate one more page when free space runs out
 - void kfree(char *addr, int size)
 - Free function
 - Memory should be allocated by kmalloc() first
 - IMPLEMENT “page deallocation”, if all slabs in each page is released
 - void slabdump()
 - Dump function for slab allocator

```

1 #include "types.h"
2 #include "defs.h"
3 #include "param.h"
4 #include "memlayout.h"
5 #include "mmu.h"
6 #include "spinlock.h"
7 #include "slab.h"
8
9 struct {
10     struct spinlock lock;
11     struct slab slab[NSLAB];
12 } stable;
13
14 void slabinit(){
15     /* fill in the blank */
16 }
17
18 char *kmalloc(int size){
19     /* fill in the blank */
20     return 0;
21 }
22
23 void kfree(char *addr, int size){
24     /* fill in the blank */
25 }
26
27 void slabdump(){
28     cprintf("____slabdump____\n");
29
30     struct slab *s;
31
32     cprintf("size\tnum_pages\tused_objects\tfree_objects\n");
33
34     for(s = stable.slab; s < &stable.slab[NSLAB]; s++){
35         cprintf("%d\t%d\t%d\t%d\n",
36                 s->size, s->num_pages, s->num_used_objects, s->num_free_objects);
37     }
38 }
```

Project 3. Byte-level Memory Allocator

- slabtest system call
 - slabtest.c, sysslab.c
 - We will test the project via various slabtest() functions
 - Please refer to the test code
 - **DO NOT MODIFY** the slabtest.c

```

/* TEST1: Single slab alloc */
cprintf("===== TEST1 =====\n");
start = counter;
t[0][0] = (int*) kmalloc (TESTSIZE);
*(t[0][0]) = counter;
counter++;

cprintf( (*t[0][0]) == start && numobj_slab(TESTSLABID) == 1) ? "OK\n": "WRONG\n";
kmfree ((char*) t[0][0], TESTSIZE);

/* TEST2: Single slab alloc: the size not equal to a power of 2. */
cprintf("===== TEST2 =====\n");
start = counter;
t[0][0] = (int*) kmalloc (TESTSIZE-10);
*(t[0][0]) = counter;
counter++;

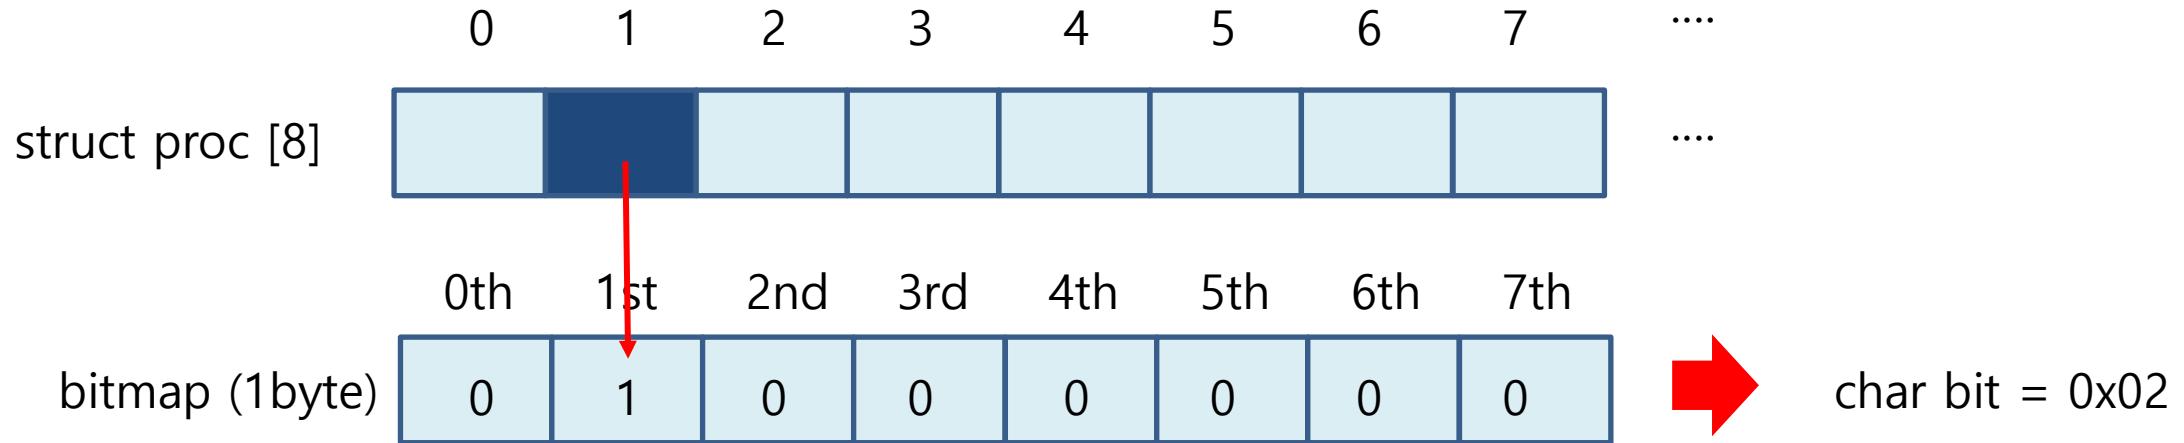
cprintf( (*t[0][0]) == start && numobj_slab(TESTSLABID) == 1) ? "OK\n": "WRONG\n";
kmfree ((char*) t[0][0], TESTSIZE);

/* TEST3: Multiple slabs alloc */
cprintf("===== TEST3 =====\n");
start = counter;
for (int i=0; i<NSLAB; i++)
{
    slabsize = 1 << (i+4);
    t[i][0] = (int*) kmalloc (slabsize);
    for (int j=0; j<slabsize/sizeof(int); j++)
    {
        memmove (t[i][0]+j, &counter, sizeof(int));
        counter++;
    }
}

// CHECK
pass = 1;
for (int i=0; i<NSLAB; i++)
{
    slabsize = 1 << (i+4);
    for (int j=0; j < slabsize/sizeof(int); j++)
    {
        //cprintf("%d, %d, %d, %d\n", i, j, *(t[i][0]+j), start);           //YOU MAY USE
        if ( *(t[i][0]+j) != start )
        {
            pass = 0;
            break;
        }
        start++;
    }
}
cprintf( pass ? "OK\n" : "WRONG\n");

```

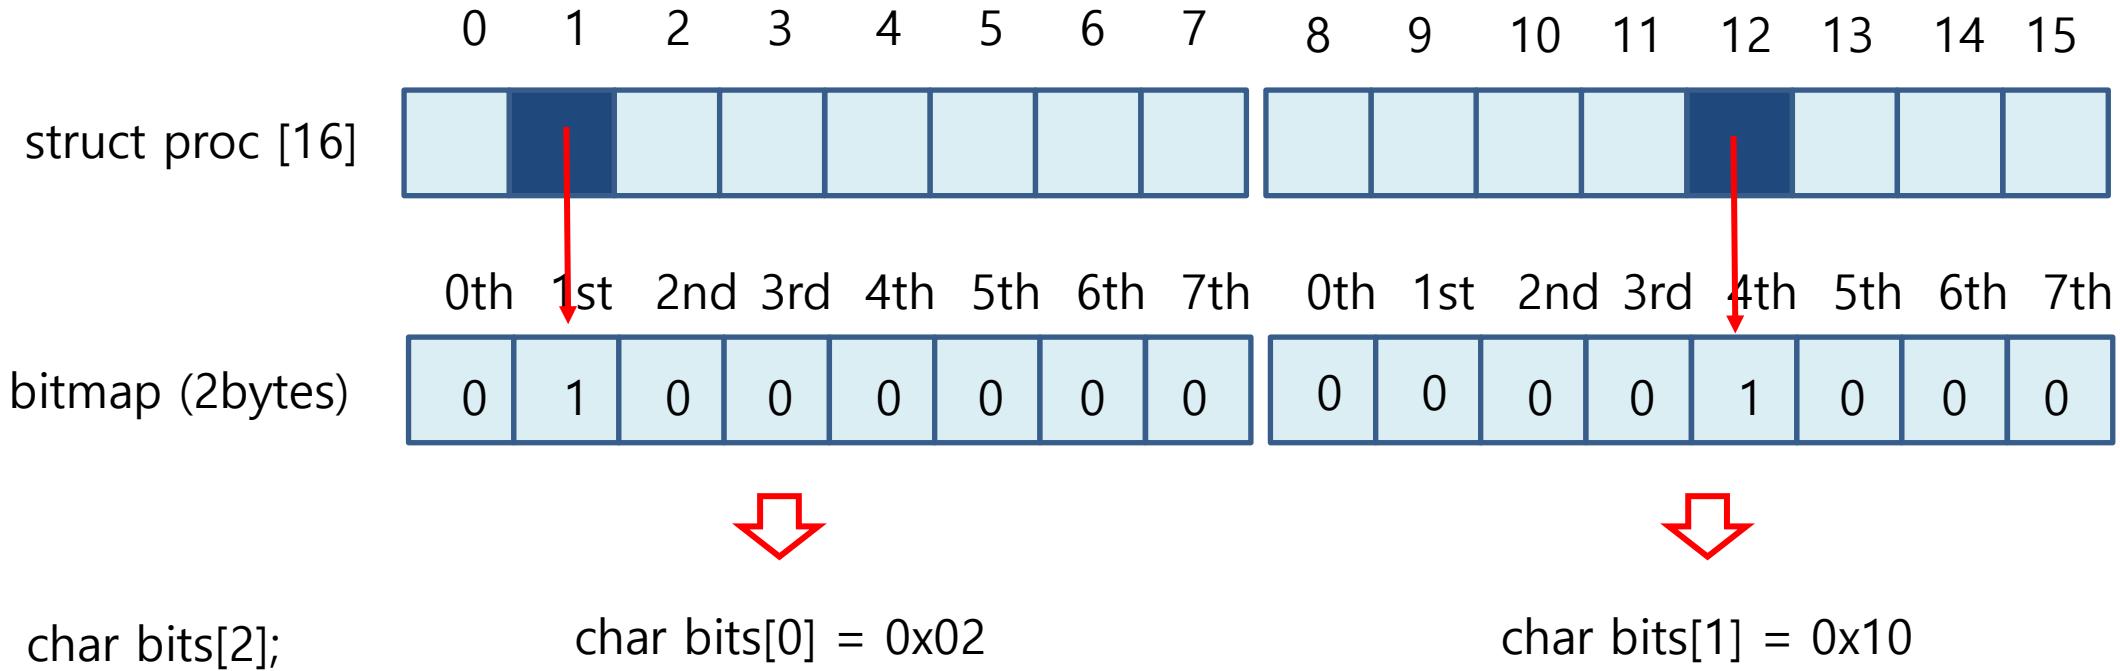
Note: bitmap



Distinguish which element is now allocating with the smallest data (i.e., a bit)

Note: In this project, use a single page (i.e., 4KB) as a bitmap for allocated slabs. Then the bitmap can distinguish $4 \times 1024 \times 8$ slabs.

Note: bitmap



Note: Bit Manipulation

```
#include<stdio.h>
#include<stdbool.h>

bool get_bit(int num, int i) {
    return ((num & (1 << i)) != 0);
}

int set_bit(int num, int i) {
    return num | (1 << i);
}

int clear_bit(int num, int i) {
    int mask = ~(1 << i);
    return num & mask;
}
```

```
int main()
{
    int num = 0;

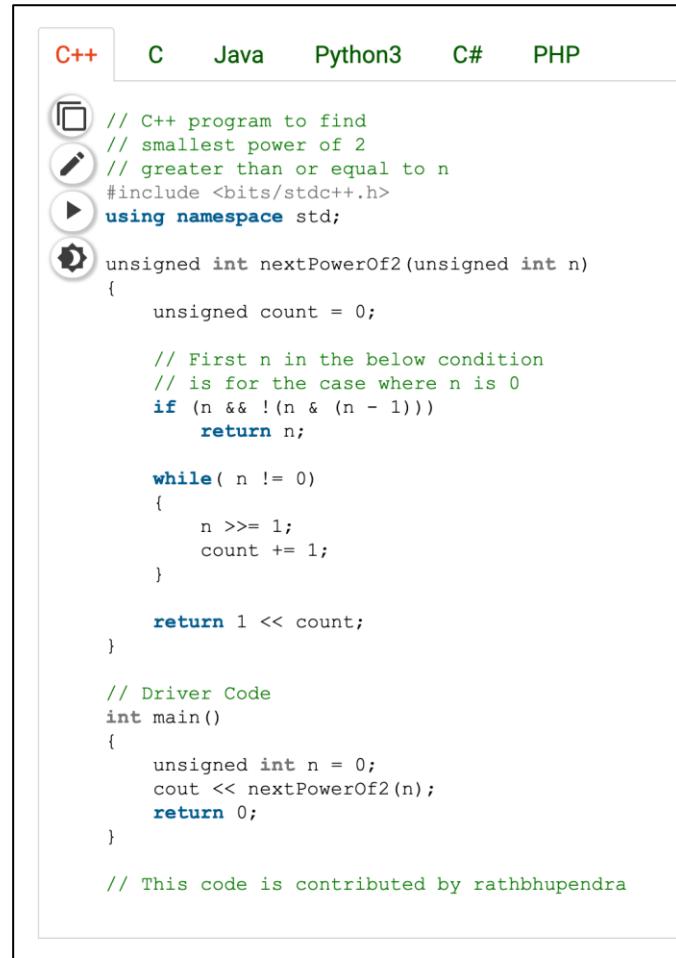
    num = set_bit(num, 1);
    printf("%x\n", num);

    if(get_bit(num, 1))
        num = clear_bit(num, 1);
    printf("%x\n", num);

}
```

Calculating Index

- Find a smallest power of 2 greater than or equal to n



The code editor interface shows a C++ program. The tabs at the top are C++, C, Java, Python3, C#, and PHP. The C++ tab is selected. The code is as follows:

```
// C++ program to find
// smallest power of 2
// greater than or equal to n
#include <bits/stdc++.h>
using namespace std;

unsigned int nextPowerOf2(unsigned int n)
{
    unsigned count = 0;

    // First n in the below condition
    // is for the case where n is 0
    if (n && !(n & (n - 1)))
        return n;

    while( n != 0)
    {
        n >= 1;
        count += 1;
    }

    return 1 << count;
}

// Driver Code
int main()
{
    unsigned int n = 0;
    cout << nextPowerOf2(n);
    return 0;
}

// This code is contributed by rathbhupendra
```

<https://www.geeksforgeeks.org/smallest-power-of-2-greater-than-or-equal-to-n/>

Project 3.

- Download template
 - <https://github.com/KilhoLee/xv6-ssu.git> (pull or clone)
 - tar xvzf xv6_ssu_slab.tar.gz
- Implement 3. slab allocator

Project 3.

- Compress your code (ID: 20201234)
 - \$tar cvzf xv6_ssu_slab_20201234.tar.gz xv6_ssu_slab
 - Please command \$make clean before compressing
- Submit your tar.gz file through **class.ssu.ac.kr**
- **PLEASE DO NOT COPY !!**

Supplements for P3

Other issues

- dynamic process allocation 구현 시 slabinit() 은 userinit() 보다 먼저 호출되어야 함.
- slabtest 수행 시 실제 xv6 가 보유한 물리 메모리보다 더 많은 양을 요구 할 경우 out-of-memory 에러나면서 중단됨.
- 과제 평가 시 총 사용 공간이 5MB 이하인 경우로 테스트 할 예정임 (함께 제공되는 slabtest.c 파일의 테스트 케이스들을 평가.)

Other issues

- 현재 xv6 kernel 에 memory linking 버그가 존재함. 특히, slabtest.c 파일에서 사용하는 string literal 에 따라, memory align이 제대로 되지 않아 커널이 실행되지 않는 경우가 발생.

```

call.o sysfile.o sysproc.o trapasm.o trap.o uart.o vectors.o vm.o sl
y initcode entryother
ld: warning: dot moved backwards before `._stab'
ld: warning: dot moved backwards before `._stab'
objdump -S kernel > kernel.asm
objdump -t kernel | sed '1,/SYMBOL TABLE/d; s/ .* / /; /^$/d' > kern...
dd if=/dev/zero of=xv6.img count=10000
10000+0 records in
10000+0 records out
5120000 bytes (5.1 MB, 4.9 MiB) copied, 0.0160489 s, 319 MB/s
dd if=bootblock of=xv6.img conv=notrunc
1+0 records in
1+0 records out
512 bytes copied, 2.5043e-05 s, 20.4 MB/s
dd if=kernel of=xv6.img seek=1 conv=notrunc
347+1 records in
347+1 records out
177924 bytes (179 kB, 174 KiB) copied, 0.000452995 s, 393 MB/s
qemu-system-i386 -nographic -drive file=fs.img,index=1,media=disk,format=raw -drive file=xv6.img,
0,media=disk,format=raw -smp 1 -m 512

```

이 메시지가 나오면 align error 가 발생한 것임.

커널이 실행되지 못하고 여기서 freezing 됨.

Other issues

- 이 경우, slabtest.c 의 string literal 길이를 조절하여 .stab section의 alignment 를 강제조절.
 - # objdump -h kernel 명령어 입력시, “.stab” section의 VMA가 4의 배수가 되어야함.

```
root@7513cfb39170:/home/khlee/xv6-ssu/solution/xv6_ssu_slab# objdump -h kernel

kernel:      file format elf32-i386

Sections:
Idx Name      Size    VMA      LMA      File off  Align
 0 .text       00006efa 80100000 00100000 00001000 2**2
               CONTENTS, ALLOC, LOAD, READONLY, CODE
 1 .rodata     00000b4b 80106f00 00106f00 00007f00 2**5
               CONTENTS, ALLOC, LOAD, READONLY, DATA
 2 .stab       00000001 80107a4b 80107a4b 00008a4b 2**0
               CONTENTS, ALLOC, LOAD, DATA
 3 .stabstr    00000001 80107a4c 80107a4c 00008a4c 2**0
               CONTENTS, ALLOC, LOAD, DATA
 4 .data       00002516 80108000 80108000 00009000 2**12
```

여기가 4의 배수가 되어야함
예를 들면, 0x80107a4c.
현재 값에서 +1을 해주면 됨.

```
1 void slabtest(){
2     int counter = 1;
3     int start;
4     int pass;
5     int* tmp;
6
7     /* NOTE:
8      * The cu
9      * especi
10     * More s
11     * to locate the VMA (virtual memory address, or an mu
12     * You can check the VMA in .stab by executing 'objdum
13     * If the VMA is not aligned to the multiple of 4B, th
14     * the literal string is in the cprintf function
15     */
16     cprintf("==== SLAB TEST ====\n");
17
18     /* TIPS:
19      * If you want to align the VMA to the multiple of 4B, you can
20      * use the alignment attribute in the .stab section. For example,
21      * alignment: 4
22      * This will ensure that the VMA is aligned to the multiple of 4B.
23      */
24 }
```

slabtest() 함수의 첫부분에 있는 이 string에 글자를 하나 더 추가 (+1) 한 후 컴파일하면, .stab VMA 위치가 0x80107a4c로 변해 4의 배수가 되어 정상 작동함.